

Can We Predict Democratic Backsliding?

Leader Traits as Early Warning Signals

Draft — Not for Circulation

March 18, 2026

Abstract

Why do some democratically elected leaders preside over limited institutional change while others oversee severe democratic backsliding? I argue that observable leader characteristics are associated with judicial erosion, but they systematically underpredict its eventual depth. Using panel data on 174 countries from 1950 to 2015, I show that irregular entry to power, military backgrounds, and political inexperience are associated with lower judicial independence. Yet across 51 episodes of severe democratic erosion, realized declines are on average fourteen times larger than trait-based predictions.

Keywords: Democratic backsliding, judicial independence, leader characteristics, prediction, early warning, Venezuela

JEL Codes: D72, K40, P48

Contents

1	Introduction	4
2	Related Literature	5
2.1	Do Leaders Matter?	5
2.2	Democratic Backsliding	6
2.3	Judicial Independence and Constitutional Compliance	8
2.4	Political Prediction	8
3	Data and Measurement	9
3.1	Judicial Independence: V-Dem	9
3.2	Leader Characteristics: Archigos and LEAD	9
3.3	Merged Panel	10
4	Empirical Strategy	10
4.1	Panel Regressions	10
4.2	Backsliding Prediction	11
4.3	Synthetic Difference-in-Differences	13
5	Results	14
5.1	Leader Traits and Judicial Constraints	14
5.2	Robustness: Alternative Outcome Variables	14
5.3	Heterogeneity by Entry Type	15
6	The Prediction Gap	17
6.1	Identifying Backsliding Episodes	17
6.2	Traits Predict Occurrence but Not Magnitude	17
6.3	Binary Classification	18
6.4	Rare Events Logit	19

6.5	Top Backsliding Episodes	19
6.6	Comparative Cases	20
7	Case Study: Venezuela	21
7.1	Background	21
7.2	Synthetic Difference-in-Differences	21
7.3	Decomposing the Venezuelan Collapse	23
8	Robustness of the Prediction Gap	24
8.1	Alternative Thresholds	24
8.2	Confidence-Interval Non-Overlap	24
8.3	Nonlinear Model Specifications	25
9	Discussion	26
9.1	Why Traits Underpredict: Type versus Agency	26
9.2	Connection to the Literature on Autocratization	27
9.3	Implications for Early Warning Systems	28
9.4	Limitations	29
10	Conclusion	30

1 Introduction

When Hugo Chávez won Venezuela’s 1998 presidential election, his biography contained clear warning signs. He was a military officer who had led a failed coup in 1992, a political outsider who had never held elected office, and a self-described revolutionary who campaigned on rewriting the constitution. Political scientists studying leader characteristics would have flagged these traits as predictive of lower judicial independence and constitutional compliance (Gutmann et al., 2024). They would have been right to flag the risk—but wrong about the magnitude. Chávez did not merely preside over a modest decline in judicial constraints. He initiated a constitutional assembly, replaced the Supreme Court, established a Judicial Emergency Committee, and systematically purged disloyal judges, driving Venezuela’s judicial independence from approximately 0.65 to below 0.05 in under a decade.

The question is: *can we predict democratic backsliding from leader traits, and if so, how well?* The answer reveals a fundamental asymmetry. Observable leader characteristics (military background, mode of entry to power, education, and political experience) are statistically significant correlates of judicial independence in a panel of 174 countries. A leader who entered power irregularly is associated with 7–11 percentage points lower judicial independence, depending on the specification. In sample, a linear probability model with country and year fixed effects correctly flags all leaders who preside over severe backsliding in the LEAD subsample (100% sensitivity at a 0.05 threshold). But the predicted magnitude of decline is, on average, $14\times$ *smaller* than what actually occurs.

I argue that this underprediction reflects a substantive point about the nature of authoritarian consolidation. Leader traits capture *type*: the kind of person who presides over declining democratic institutions. But the actual damage depends on what the leader chooses to do with power, operating through institutional channels whose dynamics traits cannot capture. A military outsider might govern moderately, or might rewrite the constitution. Observable characteristics tell us which outcome is more likely, but not by how much.

The empirical strategy proceeds in three steps. First, I estimate panel regressions of

judicial independence (from V-Dem) on leader characteristics (from Archigos and LEAD) with country and year fixed effects. Second, I identify 51 severe backsliding episodes and compare actual declines to trait-predicted declines. Third, I use Venezuela as a case study, applying synthetic difference-in-differences (Arkhangelsky et al., 2021) with Latin American comparators to estimate the causal effect of the 1999 constitutional reform.

The political economy of leaders has established that individual leaders affect national outcomes (Jones and Olken, 2005; Besley and Reynal-Querol, 2011; Gutmann et al., 2024), but the gap between statistical prediction and actual institutional damage has not been explicitly quantified for democratic backsliding. Meanwhile, the democratic erosion literature (Ginsburg and Huq, 2018; Levitsky and Ziblatt, 2018; Bermeo, 2016; Lührmann and Lindberg, 2019) has documented how backsliding operates through institutional channels, and the political prediction literature (Mueller and Rauh, 2022; Ward et al., 2010; Goldstone et al., 2010) has shown that prediction accuracy deteriorates for extreme events. By measuring exactly how far trait-based prediction falls short in the episodes that matter most, the analysis bridges these literatures. The findings have direct implications for practitioners working on democracy protection. Organizations such as the V-Dem Institute, Freedom House, and International IDEA routinely monitor democratic quality and issue early warnings. The prediction gap quantifies a fundamental limitation of leader-focused monitoring: it can identify *which* transitions warrant attention but cannot forecast *how much* damage will follow. Effective prevention therefore requires institutional safeguards, not just leader screening.

2 Related Literature

2.1 Do Leaders Matter?

A large literature examines whether individual leaders affect national outcomes. Jones and Olken (2005) exploit exogenous variation from leader deaths to show that leaders matter for economic growth, with effects concentrated in autocracies where institutional constraints are

weak. Besley and Reynal-Querol (2011) develop a model in which democratic institutions filter for leader quality, using the Archigos dataset to show that democracies select more educated leaders. Dreher et al. (2009) find that leader profession and education predict economic reform adoption.

Baturo and Elkink (2014) develop a measure of political experience (PolEx) and show that less experienced leaders are associated with worse governance outcomes. Horowitz et al. (2015) demonstrate that military experience and combat exposure shape foreign policy decisions, particularly the willingness to initiate conflict.

The closest predecessor is Gutmann et al. (2024), who use the same LEAD and Archigos datasets to show that leader traits predict constitutional compliance across 943 leaders. I differ in three ways. First, I shift the question from “do traits predict?” (yes) to “how well do they predict?” (poorly for extreme events), introducing the prediction gap as a quantitative measure of underprediction. Second, I provide a case-study estimate using SDID for Venezuela, which allows a causal decomposition of one specific collapse. Third, I evaluate binary classification performance, connecting the leader-trait literature to the early warning systems literature. The key limitation of this literature is the gap between statistical prediction and causal mechanism. Leaders with authoritarian traits may take power when institutions are already weakening. More importantly, the same traits that predict modest institutional decline cannot forecast whether a leader will pursue radical constitutional overhaul. I quantify this gap.

2.2 Democratic Backsliding

Bermeo (2016) identifies six forms of democratic backsliding, emphasizing that modern erosion occurs through “executive aggrandizement” rather than traditional coups. Levitsky and Ziblatt (2018) document how elected autocrats dismantle democratic checks through legal means: constitutional reform, court-packing, and media regulation. Ginsburg and Huq (2018) examine how constitutional design features either facilitate or impede backsliding.

Lührmann and Lindberg (2019) use V-Dem data to show that a “third wave of autocratization” is underway, distinct from earlier waves in its gradualism and legalism.

Haggard and Kaufman (2021) provide a comparative analysis of democratic decline in 77 episodes, distinguishing between backsliding driven by executive overreach and backsliding driven by military intervention. Svobik (2019) and Graham and Svobik (2020) show that partisan polarization weakens the public’s ability to serve as a democratic check, as voters tolerate norm violations by co-partisans. Geddes et al. (2014) develop a typology of autocratic regimes and show that personalist regimes, where power is concentrated in a single leader, are the most destructive to institutions.

More recent work emphasizes that modern authoritarians operate differently from their predecessors. Guriev and Treisman (2019) and Guriev and Treisman (2022) develop a theory of “informational autocrats” who maintain power through media manipulation rather than outright violence, helping explain why modern backsliding is harder to predict from traditional authoritarian profiles. Mukand and Rodrik (2020) distinguish electoral democracy from liberal democracy, showing that leaders may win elections while systematically undermining minority rights and institutional constraints.

Recent work in this journal has advanced understanding of what drives backsliding. Haggard and Tiede (2024) study the onset of judicial backsliding across 98 countries using rare events logit, showing that executive aggrandizement and political polarization predict judicial erosion. Benasaglio Berlucci and Kellam (2023) use bias-adjusted matching across 856 executives to test whether populists, presidents, or dominant executives are most responsible for democratic decline, finding that populist ideology, not institutional position, is the key driver. Both studies use V-Dem data and cross-national panels, as does this paper. I extend their work in a specific direction: while Haggard and Tiede study the *onset* of judicial backsliding and Benasaglio Berlucci and Kellam study the *average effect* of leader type, I quantify the gap between trait-predicted and actual magnitude of decline, showing that even accurate onset prediction dramatically understates the severity of backsliding episodes.

This literature shows that backsliding operates *through* institutions. It is precisely the leader’s ability to change the rules that makes backsliding dangerous and hard to predict from personal characteristics alone.

2.3 Judicial Independence and Constitutional Compliance

Judicial independence is both a primary target and a key indicator of democratic backsliding. Helmke (2002) develops a strategic theory of judicial behavior under authoritarian pressure. Ríos-Figueroa and Staton (2011) provide cross-national evidence on its determinants. Garoupa et al. (2025) use synthetic control methods to estimate the causal effect of Venezuela’s 1999 constitutional reform on judicial independence, finding a dramatic and persistent collapse.

Ginsburg (2003) shows that judicial independence emerges as a form of “constitutional insurance” when political competition is uncertain, while Stephenson (2003) models the conditions under which independent courts survive executive pressure. Voigt (2015) provides a comprehensive cross-country analysis of de facto judicial independence, distinguishing it from de jure constitutional provisions.

2.4 Political Prediction

The question of whether political events can be predicted has received growing attention. Goldstone et al. (2010) develop a model that predicts political instability with roughly 80% accuracy using a small number of variables. Ward et al. (2010) show that simple models outperform complex ones for predicting state failure. Mueller and Rauh (2022) address the “hard problem” of prediction for conflict prevention, showing that prediction accuracy deteriorates for extreme events. Acemoglu et al. (2019) provide the broader context linking institutions to growth, while Funke et al. (2023) demonstrate that populist leaders systematically damage economies, with effects that persist long after they leave office. This literature suggests that trait-based prediction of democratic backsliding may face similar accuracy

deterioration for extreme outcomes.

3 Data and Measurement

3.1 Judicial Independence: V-Dem

I use the Varieties of Democracy (V-Dem) dataset version 16, which provides expert-coded indices of institutional quality for 202 countries from 1789 to 2024 (Coppedge et al., 2024). My primary outcome variable is *judicial constraints on the executive* (hereafter “judicial independence”), which measures the extent to which the judiciary constrains executive power through review and civil liberties protection. The index ranges from 0 (no constraints) to 1 (full constraints).

A potential concern with expert-coded measures is that coders may anchor on leader characteristics when scoring institutional quality, which would mechanically inflate the estimated relationship between leader traits and judicial independence. While V-Dem’s coding protocol instructs experts to score institutional features independently of who holds office, this concern cannot be fully eliminated and should be kept in mind when interpreting the results.

As robustness checks, I examine two additional outcome variables: *high court independence* (an expert survey measure rescaled to have mean zero) and *executive respect for the constitution* (also rescaled to mean zero).

3.2 Leader Characteristics: Archigos and LEAD

Leader data come from two sources. The Archigos 4.1 dataset (Goemans et al., 2009) covers leaders of 188 countries from 1875 to 2015, providing entry and exit dates, mode of entry (regular vs. irregular), and demographics. Irregular entry includes coups, foreign installation, and unconstitutional seizures of power. I expand Archigos from leader spells to a leader-year panel.

The LEAD dataset (Horowitz et al., 2015) covers over 2,000 leaders from 1875 to 2004, providing military service, combat experience, military career, education level (0–3 scale), years of political experience, and career politician status.

3.3 Merged Panel

The full panel contains 8,975 country-leader-year observations for 174 countries. The V-Dem data extend through 2024; Archigos through 2015; and LEAD through 2004. Regression samples are slightly smaller (8,961 in the Archigos specification) because observations with missing leader age are excluded. In the raw data, 45 leader-year observations have implausible age values (below 0 or above 110) and are set to missing; of these, 13 survive into the analysis sample after restricting to post-1950 country-years with non-missing judicial independence. One additional observation is dropped due to missing birth year, yielding $N = 8,961$. The LEAD subsample covers 1,143 observations for 155 countries through 2004. Table 1 defines all variables and Table 2 presents summary statistics. The variation in N across LEAD variables (1,194–1,262) reflects differential coverage of specific traits rather than sample restrictions; the listwise-complete LEAD subsample used in the main regressions contains 1,143 observations.

4 Empirical Strategy

4.1 Panel Regressions

I estimate:

$$Y_{it} = \beta' \mathbf{X}_{it} + \alpha_i + \gamma_t + \varepsilon_{it} \quad (1)$$

where Y_{it} is judicial independence for country i in year t , $\mathbf{X}_{it}$ is a vector of leader characteristics, α_i are country fixed effects, γ_t are year fixed effects, ε_{it} is the error term, and standard errors are clustered by country. Variation comes from within-country changes:

when a country’s leader changes type, how does institutional quality change relative to the global trend?

The key identification concern is that institutional weakness may cause the selection of authoritarian-type leaders, creating reverse causality. Countries with eroding institutions may be more likely to experience irregular leader entry, which would bias the estimated coefficients upward. Country fixed effects address time-invariant confounders but cannot resolve time-varying reverse causality. The estimates should therefore be interpreted as within-country associations rather than causal effects. I return to this concern in the Limitations section.

I present three specifications in a stepwise fashion: (1) Archigos traits only (irregular entry, female, leader age), which exploits the full sample; (2) adding LEAD military traits (military service, combat, military career); and (3) the full set of LEAD traits including education, political experience, and career politician status.

4.2 Backsliding Prediction

Using fitted values from the panel regression, I compute predicted institutional quality for each country-year. I then construct a binary backsliding indicator and estimate a linear probability model (LPM):

$$B_{it} = \boldsymbol{\pi}'\mathbf{X}_{it} + \alpha_i + \gamma_t + u_{it} \quad (2)$$

where $B_{it} = 1$ if the country experienced a severe decline in judicial independence over the preceding five years, $\mathbf{X}_{it}$ is the same vector of leader characteristics, α_i and γ_t are country and year fixed effects, u_{it} is the error term, and $\boldsymbol{\pi}$ is the LPM coefficient vector. Standard errors are clustered by country. The LPM is preferred over logit here because the inclusion of high-dimensional fixed effects makes maximum likelihood estimation of nonlinear models unreliable (Greene, 2004). I evaluate in-sample classification accuracy using sensitivity (true positive rate) and specificity (true negative rate) at a 0.05 probability threshold. This

threshold is deliberately low, chosen to maximize sensitivity for an early warning application where missing a true backslider is more costly than generating false alarms.

A potential concern with the LPM is that backsliding episodes are rare events, constituting approximately 2.3% of country-year observations. When the dependent variable is this imbalanced, standard logistic regression can produce biased coefficient estimates and underestimate the probability of the rare outcome in finite samples (King and Zeng, 2001). The bias arises because maximum likelihood estimation of logit models is consistent but not unbiased; in small samples or with rare outcomes, the estimates are systematically shifted toward zero, attenuating both coefficients and predicted probabilities. King and Zeng (2001) propose a correction that applies a bias-reduction penalty to the likelihood function and adjusts predicted probabilities using prior information about the population event rate. This rare events logit estimator has been widely adopted in conflict studies (Beck et al., 2000), international relations, and, most relevant to this paper, the judicial backsliding literature (Haggard and Tiede, 2024). I therefore re-estimate the binary classification model using rare events logit as a robustness check (Section 6.4). Because the King and Zeng (2001) estimator cannot accommodate high-dimensional fixed effects, the rare events logit specification replaces country and year fixed effects with decade dummies, providing a complementary test of whether leader traits predict backsliding even without absorbing country-specific heterogeneity.

I define a backsliding episode as a 5-year decline in judicial constraints on the executive exceeding 0.20 (roughly two-thirds of a standard deviation). This threshold is conservative relative to the existing literature. Lührmann and Lindberg (2019) define autocratization episodes as a cumulative decline of 0.10 or more on the Electoral Democracy Index, while the V-Dem Institute’s annual democracy reports use the same 0.10 threshold for “substantial autocratization.” My 0.20 threshold on a narrower indicator (judicial constraints rather than electoral democracy) is designed to identify only *severe* episodes of institutional erosion, reducing the risk of false positives from measurement noise or minor political fluctuations. This

conservative approach follows the recommendation of Haggard and Kaufman (2021), who argue for restricting attention to cases where democratic decline is statistically significant and substantively meaningful. Section 8 demonstrates that results are robust to alternative backsliding definitions, including thresholds of 0.15 and 0.25, and a confidence-interval non-overlap method following Haggard and Tiede (2024).

For each backsliding episode, I compare the actual 5-year decline to the trait-predicted decline (computed from the fitted values of Equation 1). The ratio of actual to predicted decline provides a measure of the “prediction gap.”

4.3 Synthetic Difference-in-Differences

For the Venezuela case study, I use the synthetic difference-in-differences (SDID) estimator of Arkhangelsky et al. (2021). SDID generalizes both standard difference-in-differences and the synthetic control method of Abadie et al. (2010). The estimator solves:

$$(\hat{\tau}^{\text{sdid}}, \hat{\mu}, \hat{\alpha}, \hat{\gamma}) = \arg \min_{\tau, \mu, \alpha, \gamma} \left\{ \sum_{i=1}^N \sum_{t=1}^T (Y_{it} - \mu - \alpha_i - \gamma_t - W_{it}\tau)^2 \times \hat{\omega}_i^{\text{sdid}} \hat{\lambda}_t^{\text{sdid}} \right\} \quad (3)$$

where Y_{it} is judicial independence for country i in year t , W_{it} is the treatment indicator (equal to 1 for Venezuela after 1999), μ is an intercept, α_i and γ_t are unit and time fixed effects (using the same notation as in Equations 1–2, though here they are jointly estimated with the SDID weights rather than absorbed via demeaning), and $\hat{\omega}_i^{\text{sdid}}$ and $\hat{\lambda}_t^{\text{sdid}}$ are unit and time weights chosen to improve pre-treatment fit. The unit weights $\hat{\omega}_i$ reweight the donor pool (17 Latin American countries) so that their weighted average tracks Venezuela’s pre-treatment trajectory, as in the synthetic control method. The time weights $\hat{\lambda}_t$ downweight pre-treatment periods that are less informative about the treatment effect, as in a DiD framework. Standard errors are obtained via placebo inference (50 replications), assigning treatment status to each donor unit in turn and computing the distribution of placebo effects.

The small number of donors (17) limits the resolution of placebo inference; the resulting p-values should be interpreted as approximate.

This approach avoids the sensitivity to predictor selection that affects traditional synthetic control (Abadie et al., 2010) and provides valid inference with a single treated unit. Compared to standard DiD, SDID relaxes the parallel trends assumption by reweighting both units and time periods to improve pre-treatment balance.

5 Results

5.1 Leader Traits and Judicial Constraints

Table 3 presents results for judicial constraints on the executive, the primary outcome. The most consistent finding across specifications is that *irregular entry to power* is associated with lower judicial independence: -0.106 in the full Archigos sample ($p < 0.001$, $N = 8,961$), declining to -0.075 in the smaller LEAD sample ($p < 0.001$, $N = 1,143$). Military career is also negatively associated (-0.038 , $p < 0.10$). Leader age is negatively associated with judicial independence across all specifications, with the coefficient becoming larger and more precisely estimated as additional controls are added (-0.001 to -0.002). The within- R^2 ranges from 0.066 in the Archigos-only specification to 0.122 in the full LEAD model, indicating that leader traits explain roughly 7–12% of within-country variation in judicial independence, leaving the vast majority unexplained.

5.2 Robustness: Alternative Outcome Variables

Tables 4 and 5 replicate the analysis using high court independence and executive respect for the constitution. The results are broadly consistent. Irregular entry predicts lower scores on all three measures: -0.402 ($p < 0.01$) for high court independence and -0.652 ($p < 0.01$) for executive respect in the full Archigos sample. Military service is a strong predictor of lower constitutional respect (-0.291 , $p < 0.05$) but not of high court independence. Education

is positively associated with high court independence (+0.087, $p < 0.05$), consistent with the Besley and Reynal-Querol (2011) finding that democracies select more educated leaders. These results suggest that the association between leader traits and institutional quality is not an artifact of any single measure.

5.3 Heterogeneity by Entry Type

The main results pool leaders who entered power through elections with those who seized power through coups or revolutions. If the mechanisms linking personal traits to institutional outcomes differ across these two paths to power, the pooled estimates may obscure important variation. I investigate this by estimating the full LEAD specification separately for the regular-entry subsample ($N = 920$) and the irregular-entry subsample ($N = 157$). Table 6 reports the split-sample results.

The most striking difference concerns political experience. Years of prior political experience is near zero and statistically insignificant among elected leaders ($\hat{\beta} = 0.000$, $p = 0.59$) but large and highly significant among leaders who seized power irregularly ($\hat{\beta} = 0.006$, $p = 0.01$ for judicial constraints; $\hat{\beta} = 0.037$, $p < 0.01$ for high court independence; $\hat{\beta} = 0.042$, $p = 0.01$ for executive respect). This correlation is consistent with experienced coup leaders presiding over *higher* judicial independence, though the small subsample ($N = 157$) and the possibility of reverse causality warrant caution in interpretation.

This pattern is consistent with a growing literature on institutional variation within authoritarian regimes. Svolik (2012) argues that dictators who must share power with a ruling coalition—particularly military juntas—maintain institutions such as legislatures and courts as credible commitment devices to prevent coups by their own allies. Gandhi and Przeworski (2007) show that autocrats who face organized opposition maintain nominally democratic institutions to co-opt potential challengers, and that this strategy extends their tenure. In both frameworks, politically experienced leaders who seize power are more likely to understand and deploy institutional mechanisms, while inexperienced outsiders are more

likely to resort to pure coercion.

The historical record offers clear illustrations. In Chile, Pinochet—a career military officer with decades of institutional experience—maintained some formal judicial structures and eventually submitted to a plebiscite, facilitating a negotiated transition (Barros, 2002). In contrast, Idi Amin in Uganda, a military officer with no political career who seized power in 1971, immediately suspended the constitution and dismantled the courts. More broadly, Geddes et al. (2014) document that military regimes led by officers embedded in hierarchical command structures are far more likely to negotiate institutional exits than personalist dictators, with nearly 62% of military regimes eventually democratizing compared to only 36% of personalist regimes. Among leaders who enter through irregular channels, political experience thus proxies for the leader’s orientation toward institutional governance versus pure coercion—a distinction that is irrelevant among elected leaders, where experience varies within a narrower band of career politicians.

Education level also shifts across subsamples: it is positively associated with judicial independence among elected leaders ($\hat{\beta} = 0.011$, $p < 0.10$) but negatively associated among irregular-entry leaders ($\hat{\beta} = -0.038$, $p < 0.10$). The career-politician indicator is insignificant for elected leaders but negative and marginally significant for irregular-entry leaders ($\hat{\beta} = -0.087$, $p < 0.10$).

I test these differences formally using a fully interacted model on the combined LEAD sample ($N = 1,143$), interacting irregular entry with every leader trait. Joint F -tests for the null that all interaction terms equal zero yield $F = 0.70$ ($p = 0.69$) for judicial constraints, $F = 1.03$ ($p = 0.42$) for high court independence, and $F = 1.31$ ($p = 0.24$) for executive respect for the constitution. I cannot reject that the trait-effect vector operates identically across entry types as a whole, though individual coefficients—particularly political experience—shift meaningfully.

6 The Prediction Gap

6.1 Identifying Backsliding Episodes

Applying the 0.20 threshold to 5-year changes in judicial constraints, I identify 51 distinct backsliding episodes across 184 country-years (182 with non-missing Archigos predictions). These episodes span a wide range of contexts: military coups (Pinochet’s Chile in 1974, Ne Win’s Burma in 1962), revolutionary takeovers (Castro’s Cuba in 1959), post-Cold War authoritarian consolidation (Lukashenko’s Belarus in 1996), and executive aggrandizement through legal reform (Chávez’s Venezuela in 2000). The diversity of these cases suggests that the prediction gap is not driven by a single type of backsliding.

6.2 Traits Predict Occurrence but Not Magnitude

Using fitted values from the Archigos panel regression (Column 1 of Table 3), I compute the predicted 5-year change in judicial independence for each country-year. Figure 1 plots actual against predicted changes. For non-backsliding observations (gray), actual and predicted changes cluster near zero and along the 45-degree line. For backsliding episodes (red), actual declines are far larger than predicted, clustering in the lower portion of the plot.

Table 7 quantifies the gap. The average backsliding episode involves a 0.361-point decline, but leader traits predict only a 0.026-point decline. The underprediction ratio is $14.0\times$. This pattern is not driven by outliers; it holds across all 51 episodes. To benchmark this ratio, I conduct a Monte Carlo simulation (1,000 replications) generating data from the estimated model under homogeneous normal residuals. Under this mechanical benchmark, the expected underprediction ratio is approximately $49\times$, reflecting the fact that any linear model will severely underpredict tail events driven by idiosyncratic shocks. The actual ratio of $14\times$ is substantially *lower* than this benchmark, indicating that leader traits capture real predictive content, reducing the mechanical underprediction by roughly 70%. Nevertheless, the remaining $14\times$ gap is substantively large: traits help, but they leave the vast majority

of actual backsliding unexplained.

6.3 Binary Classification

Despite the magnitude gap, leader traits perform well as in-sample binary classifiers. Table 8 reports the LPM estimates. A linear probability model (Equation 2) of backsliding occurrence achieves 100% in-sample sensitivity (correctly flagging all 28 episodes in the LEAD sample) and 82% specificity (18% false positive rate) at a threshold of 0.05. Irregular entry is the strongest predictor (+3.3 percentage points, $p < 0.01$ in the full Archigos sample). In the LEAD subsample, combat experience is negatively associated with backsliding (−4.9 percentage points, $p < 0.05$), suggesting that leaders with direct military experience may be less likely to undermine institutions than those with purely political military careers.

These classification results should be interpreted with caution. Because the LPM includes high-dimensional country and year fixed effects, the model has substantial flexibility to fit the training data. As a preliminary out-of-sample check, I estimate the LPM on pre-1990 data and predict post-1990 backsliding episodes. The out-of-sample sensitivity is 94.4% (68 of 72 post-1990 backsliding country-years correctly flagged), with specificity of 83.8%. These results suggest that the in-sample classification performance is not purely an artifact of overfitting, though the temporal split does not address the concern that fixed effects absorb much of the predictive variation.

The 18% false positive rate warrants careful interpretation. As Figure 1 reveals, many country-years with negative predicted changes do not experience actual backsliding. These false positives arise because the model identifies the *type* of leader associated with institutional decline, yet many such leaders govern without dismantling the judiciary. This asymmetry is informative: it shows that having “risky” traits is a necessary but not sufficient condition for severe institutional erosion. The false positive rate is also consistent with the broader early warning literature. Ward et al. (2010) find similar tradeoffs in conflict prediction models, where high sensitivity comes at the cost of over-flagging stable cases.

The asymmetry is clear: traits identify *who might* backslide with reasonable accuracy, but they cannot predict *how far*.

6.4 Rare Events Logit

Backsliding episodes constitute approximately 2.3% of country-year observations (184 out of 8,106). Standard logistic regression can produce biased coefficient estimates and predicted probabilities when the dependent variable is this rare (King and Zeng, 2001). Following Haggard and Tiede (2024), who use rare events logit to predict judicial backsliding onset, I re-estimate the binary prediction model using the King and Zeng (2001) correction.

Table 9 reports the results. Because the rare events logit estimator cannot accommodate the high-dimensional country fixed effects used in the main LPM specification, both models here include decade fixed effects (with the 2010s as the reference category) instead of country and year fixed effects. This is a more demanding test of whether leader traits predict backsliding even without controlling for country-specific factors.

The results are virtually identical across the two estimators. The coefficient on irregular entry is 1.316 (SE = 0.152) in the standard logit and 1.314 (SE = 0.152) in the rare events logit. Both models yield 40.4% sensitivity and approximately 90% specificity at a 0.05 probability threshold. The rare events correction has negligible impact in this application, likely because the sample size ($N = 8,080$) is large enough that the finite-sample bias corrected by King and Zeng (2001) is small. This confirms that the classification results are not driven by bias from the rarity of backsliding events.

6.5 Top Backsliding Episodes

Figure 2 displays the ten most severe episodes, comparing actual to trait-predicted magnitude. The episodes shown are Belarus under Lukashenko (1996), the Philippines under Marcos (1972), Republic of the Congo under Debat (1968), Chile under Pinochet (1974), Zambia under Kaunda (1972), Cuba under Castro (1959), Venezuela under Chávez (2000),

Comoros (1975), Peru under Fujimori (1993), and Turkey under Erdogan (2016). In every case, the actual decline dwarfs the prediction. Belarus experienced a 0.53-point decline in judicial constraints; the Philippines a 0.53-point decline; Chile a 0.43-point decline. Predicted declines are uniformly small, ranging from near zero to 0.13 points.

6.6 Comparative Cases

The prediction gap is not confined to any single episode. Two contrasting cases from Figure 2 illustrate that the pattern holds regardless of how a leader enters power.

Turkey (Erdogan, 2016): Regular Entry. Recep Tayyip Erdogan entered power through regular elections in 2003, with substantial political experience as a former mayor of Istanbul. His observable traits — democratic entry, no military background, prior governing experience — predicted only a modest decline in judicial independence. Yet following the 2016 coup attempt, Erdogan pursued sweeping judicial purges, dismissing thousands of judges and prosecutors under emergency decree powers, and a 2017 constitutional referendum that subordinated the judiciary to executive appointment. Turkey’s judicial constraints index fell by over 0.40 points. (Because Erdogan’s tenure extends beyond the LEAD dataset’s 2004 endpoint, this case uses Archigos variables only.) Erdogan’s case represents the “informational autocrat” pathway described by Guriev and Treisman (2019)—gradual institutional erosion through formally legal channels—precisely the mechanism that leader traits are least equipped to predict. Here, traits *failed to flag any risk at all*.

Chile (Pinochet, 1974): Irregular Entry. Augusto Pinochet seized power through a military coup in September 1973, overthrowing the democratically elected government of Salvador Allende. His traits — irregular entry via coup, military career, no prior civilian political experience — are precisely those the prediction model associates with democratic erosion. In this sense, Pinochet’s traits correctly identified him as a *high-risk leader*. Yet even here, the model dramatically underpredicted the magnitude: Chile’s judicial constraints index fell by 0.43 points as Pinochet dissolved Congress, banned political parties, and subor-

minated the judiciary to military tribunals. The predicted decline was a fraction of the actual collapse. Pinochet thus illustrates the complementary failure: even when traits successfully flag risk, they cannot forecast how far the institutional destruction extends.

Together, these cases reveal that the prediction gap operates through two distinct channels. For leaders who enter through regular channels (Erdogan), traits fail at both detection and magnitude — the model neither flags the risk nor anticipates the severity. For leaders who enter irregularly (Pinochet), traits succeed at detection but fail at magnitude — the model identifies the risk but vastly underestimates the scale. In both cases, the gap reflects actions taken *after* assuming power that are not predictable from pre-office characteristics. Both Haggard and Tiede (2024) and Benasaglio Berlucchi and Kellam (2023) document similar patterns, reinforcing that this finding is systematic rather than idiosyncratic.

7 Case Study: Venezuela

7.1 Background

Venezuela illustrates the prediction gap in its starkest form. In February 1999, Chávez took office after winning a democratic election. Within months, he initiated a constitutional assembly that produced the 1999 Bolivarian Constitution. The new constitution restructured the judiciary, created a Judicial Emergency Committee with the power to dismiss judges without due process, and expanded the Supreme Tribunal from 15 to 32 members. Over the following decade, Chávez packed the court with loyalists and systematically purged disloyal judges (Garoupa et al., 2025).

7.2 Synthetic Difference-in-Differences

I estimate the effect of Venezuela’s 1999 reform using the SDID estimator (Equation 3), with 17 Latin American countries as the donor pool and the period 1985–2015. The pre-treatment window begins in 1985 to focus on the period of Latin American democratization

most relevant to Venezuela’s institutional trajectory; results are robust to extending the pre-treatment period to 1960 (Table 10, Panel D). The unit weights $\hat{\omega}_i$ construct a synthetic Venezuela that matches the pre-1999 trajectory of judicial independence, while the time weights $\hat{\lambda}_t$ downweight pre-treatment periods that are less informative about the treatment effect.

Figure 3 shows Venezuela’s trajectory against the SDID synthetic control. In the pre-treatment period (1985–1998), both series exhibit stable, roughly parallel trajectories: Venezuela fluctuates between 0.65 and 0.75, while the synthetic control remains between 0.53 and 0.57. The persistent level difference is absorbed by the unit fixed effect (α_i) in the SDID estimator; what matters for identification is that the *trends* are parallel, not the levels. After 1999, the two series diverge sharply: Venezuela’s judicial independence collapses from approximately 0.58 to below 0.05, while the synthetic control remains stable in the 0.60–0.65 range.

Table 10 reports the main estimate and robustness checks. The SDID estimate is -0.633 ($SE = 0.215$, $t = -2.94$), indicating that Venezuela’s judicial independence fell by approximately 0.63 points relative to its synthetic counterpart. This estimate is robust to the choice of method: standard DiD yields -0.650 and pure synthetic control yields -0.599 . Leave-one-out tests (excluding individual donors) produce estimates ranging from -0.626 to -0.659 , indicating that no single donor country drives the result. An in-time placebo test, assigning a false treatment date of 1992, yields an essentially zero effect (-0.011 , $t = -0.07$), confirming that the post-1999 divergence is not a continuation of pre-existing trends. Extending the pre-treatment window to 1960 yields a larger estimate of -0.698 (Panel D), consistent with the main finding.

The stability of the SDID estimate across specifications is reassuring. The narrow range of leave-one-out estimates (-0.626 to -0.659) indicates that the synthetic control is not dominated by any single donor. The near-zero in-time placebo provides direct evidence against the alternative hypothesis that Venezuela was already diverging from its neighbors before Chávez took office. Taken together, these results are consistent with a causal inter-

pretation: the 1999 Bolivarian Constitution, and the institutional changes it enabled, are estimated to have contributed to a large and persistent decline in judicial independence. The single-unit design means that this interpretation rests on the assumption that no other contemporaneous shock (beyond the constitutional reform and the broader Chávez political project) drove the decline.

7.3 Decomposing the Venezuelan Collapse

As an illustrative exercise, I decompose the Venezuelan gap using the panel regression coefficients. Because the panel coefficients are correlational (reflecting average within-country associations, not causal effects), this decomposition should be interpreted as a descriptive accounting exercise rather than a causal decomposition. Let $\Delta\mathbf{X} = \mathbf{X}_{\text{Chávez}} - \mathbf{X}_{\text{Caldera}}$ denote the vector of trait changes from Rafael Caldera to Hugo Chávez. The decomposition is:

$$\hat{\tau}^{\text{sdid}} = \underbrace{\hat{\boldsymbol{\beta}}' \Delta\mathbf{X}}_{\text{Trait-predicted}} + \underbrace{\hat{\tau}^{\text{sdid}} - \hat{\boldsymbol{\beta}}' \Delta\mathbf{X}}_{\text{Prediction gap}} \quad (4)$$

Table 11 applies Equation (4) to the trait changes from Caldera to Chávez. Caldera was a civilian career politician with a law degree and 30 years of experience. Chávez was a military officer, a former coup leader, and a political newcomer with only 5 years of experience. Including the age difference (Caldera was 78, Chávez was 45), the net trait-predicted change is essentially zero (+0.005): the age effect offsets the military and experience variables. The entire 0.633-point SDID estimate therefore falls into the prediction gap: none of the Venezuelan judicial collapse is explained by the change in observable leader traits. This residual reflects some combination of Chávez’s specific choices (rewriting the constitution, packing the courts, purging the judiciary), omitted contextual variables (Venezuela’s oil wealth, political polarization, institutional fragility), and model misspecification.

8 Robustness of the Prediction Gap

The baseline analysis defines backsliding episodes using an arbitrary threshold (5-year decline in judicial constraints exceeding 0.20). This section demonstrates that the prediction gap is robust to alternative definitions.

8.1 Alternative Thresholds

Table 12 reports the prediction gap under three threshold definitions (0.15, 0.20, and 0.25) and a confidence-interval non-overlap method. Lowering the threshold to 0.15 captures 290 country-years and yields a ratio of $14.2\times$. Raising it to 0.25 restricts attention to 141 country-years and yields $20.8\times$. The prediction gap is not an artifact of the baseline threshold; if anything, it grows larger for more severe episodes.

8.2 Confidence-Interval Non-Overlap

Following Haggard and Tiede (2024), I implement a statistically grounded alternative. V-Dem provides 95% credible intervals for each country-year score. I define a backsliding episode as occurring when the upper bound of the current 95% credible interval falls below the lower bound of the credible interval in any of the preceding five years. This ensures that only *statistically significant* declines are flagged, rather than relying on a researcher-chosen magnitude threshold.

The CI non-overlap method identifies 55 distinct backsliding episodes across 197 country-years. There is substantial overlap with the threshold-based definition: 150 country-years (81.5% of the threshold-based episodes) are also flagged by the CI method, but the CI method additionally identifies 40 country-years that fall below the 0.20 threshold yet represent statistically significant declines. The mean actual 5-year drop during CI-overlap episodes is -0.339 (compared to -0.361 under the threshold definition), and the mean predicted drop is -0.020 . The underprediction ratio is $16.7\times$, slightly larger than the baseline $14.0\times$.

Binary classification using the CI-overlap definition yields 94.4% sensitivity and 81.9% specificity, comparable to the baseline results. Figure 4 replicates the prediction scatter plot using the CI-overlap definition, confirming that the visual pattern of massive underprediction is robust.

These results provide strong evidence that the prediction gap is not driven by the choice of backsliding definition. Whether backsliding is defined by a fixed threshold or by statistical significance of the decline, leader traits predict occurrence but dramatically underpredict magnitude.

8.3 Nonlinear Model Specifications

A natural concern is that the underprediction reflects the linearity of OLS rather than a genuine feature of backsliding dynamics. To address this, I re-estimate the prediction gap using four alternative model specifications: quantile regression at $\tau = 0.10$, 0.25, and 0.50 (which directly model conditional quantiles rather than the mean); and two machine learning methods—random forest and gradient boosting—estimated via leave-one-country-out cross-validation (LOCO-CV) to provide honest out-of-sample predictions.

Table 13 reports the results. Quantile regressions, despite being designed to capture tail behavior, produce *larger* prediction gaps (21–36 $\times$) than OLS (13.6 $\times$). The machine learning models perform best: random forest achieves 4.1 $\times$ and gradient boosting 4.6 $\times$, reducing the gap substantially relative to OLS. However, even these flexible nonlinear models underpredict actual backsliding by a factor of four to five. The prediction gap is not an artifact of linear functional form; it persists across the full range of model specifications, from parametric to nonparametric, from in-sample to out-of-sample.

9 Discussion

9.1 Why Traits Underpredict: Type versus Agency

The $14\times$ underprediction ratio reflects a distinction between what I call *type* and *agency*. Leader traits capture type: the statistical tendency of military-background, irregular-entry leaders to preside over lower judicial independence. Panel regression coefficients are estimated from the *average* within-country effect of leader transitions, including the many cases where military leaders govern without dismantling courts. The coefficient on irregular entry (-0.075 in the LEAD subsample, -0.106 in the full Archigos sample) reflects the average transition, not worst cases.

Actual backsliding episodes involve extreme outcomes: constitutional overhauls, emergency powers, court-packing, and systematic purges. These are actions that go far beyond what observable traits predict. The prediction gap provides an upper-bound measure of how much a leader’s actions exceeded what their biography would forecast. I describe this as “agency,” with the important caveat that the residual also captures omitted variables (such as economic crises, political polarization, and pre-existing institutional vulnerabilities) and model misspecification that may amplify or attenuate the true leader-specific component.

This finding connects to Geddes et al. (2014)’s distinction between different autocratic regime types. Personalist regimes, where a single leader concentrates power by dismantling all institutional checks, produce the most severe backsliding. But “personalism” is itself an outcome of leader choices, not an observable *ex ante* trait. The traits that predict a leader’s *tendency* toward personalism (military background, irregular entry) cannot predict whether personalist consolidation will actually occur, because it depends on the leader’s strategic decisions and on context-specific institutional vulnerabilities.

9.2 Connection to the Literature on Autocratization

The prediction gap concept speaks directly to debates about what drives democratic backsliding. Bermeo (2016) documents that modern autocratization increasingly occurs through “executive aggrandizement” rather than outright coups, making it harder to detect from leader backgrounds alone. Levitsky and Ziblatt (2018) emphasize that the most effective authoritarian leaders are those who use democratic institutions against democracy itself. The results here provide quantitative support for this view: the institutional mechanisms through which backsliding operates (constitutional reform, judicial restructuring) are precisely the channels that leader traits cannot capture.

Lührmann and Lindberg (2019) show that the current “third wave of autocratization” is characterized by gradualism and legalism. This is consistent with the finding that traits predict backsliding occurrence but not magnitude: the initial steps toward autocratization may correlate with leader type, but the subsequent escalation depends on strategic choices, institutional vulnerabilities, and political opportunities that are inherently difficult to forecast.

The prediction gap also speaks to the role of populist ideology. Benasaglio Berlucchi and Kellam (2023) find that populism, not presidentialism or supermajority control, is the strongest predictor of democratic backsliding. The leader trait variables used here (from Archigos and LEAD) do not capture ideology or populist worldview; they measure career backgrounds and modes of entry. This is both a limitation and an insight. The prediction gap likely reflects, in part, the ideological dimension that observable traits miss: a military outsider with populist convictions may pursue radical institutional overhaul, while one without such convictions may govern moderately. However, even adding populism measures would be unlikely to close the gap entirely, because the gap captures actions taken *after* assuming office, not pre-existing characteristics of any kind. The distinction between trait-based prediction and action-based outcomes is fundamental, not merely a matter of omitted variables.

The results also relate to Svobik (2019)’s finding that democratic erosion is easier when the public is polarized. Leader traits may interact with social conditions in ways that amplify or attenuate their effects. A military outsider in a polarized society may find it easier to dismantle institutions than the same leader in a more cohesive polity. This suggests that prediction models combining leader traits with contextual factors may outperform trait-only models, though this extension is left for future work.

9.3 Implications for Early Warning Systems

The results suggest a two-stage approach to democratic early warning, with the caveat that these implications rest on within-sample findings and may not generalize directly to contemporary cases:

1. **Screening (feasible from traits):** Use leader characteristics to flag high-risk transitions. In-sample binary prediction achieves 100% sensitivity with 82% specificity in the LEAD subsample (28 episodes). These in-sample results would need out-of-sample validation before informing diplomatic attention and civil society monitoring.
2. **Monitoring (requires institutional observation):** Predict the severity of backsliding by tracking institutional actions: constitutional amendments, judicial appointments, emergency powers, and media regulation. This stage cannot rely on leader characteristics.

If leader traits could predict both occurrence *and* magnitude, the remedy would be candidate screening. But since traits only predict occurrence, the primary defense must be institutional: entrenchment of judicial independence, supermajority requirements for constitutional change, and independent appointment mechanisms that constrain leaders regardless of their type. For practitioners, the two-stage framework translates directly into resource allocation. Stage 1 (screening) is inexpensive and can be automated using publicly available leader biography data. Stage 2 (monitoring) is resource-intensive but can be targeted: when

a high-risk leader takes power, democracy monitors should escalate institutional tracking, including judicial appointment patterns, emergency powers invocations, and constitutional amendment proposals. The prediction gap provides a quantitative benchmark for calibrating the intensity of monitoring required.

9.4 Limitations

This analysis has three main limitations. First, the panel regressions identify leader-trait associations from within-country variation, but this variation is not exogenous. Weak institutions may produce the conditions for irregular leader entry, creating reverse causality that country fixed effects cannot resolve. V-Dem expert-coded indices may also be subject to anchoring bias if coders rate institutional quality lower for leaders with known authoritarian backgrounds. The analysis further tests 9 traits across 3 outcomes without correction for multiple hypothesis testing; the robustness of irregular entry mitigates this concern for the main finding, but weaker results should be interpreted cautiously.

Second, the LEAD dataset ends in 2004, excluding recent cases of backsliding in Hungary, Poland, and Turkey. The LEAD subsample contains only $\sim 1,000$ observations, and the SDID case study for Venezuela relies on parallel trends assumptions that cannot be definitively confirmed despite a near-zero in-time placebo estimate (-0.011 , $t = -0.07$). The small donor pool (17 countries) further limits the power of placebo inference.

Third, the prediction gap reflects not only leader agency but also omitted contextual variables (institutional vulnerability, political polarization) and model limitations. Moreover, backsliding episodes are identified post hoc from realized declines; leaders who attempted erosion but were constrained by effective institutions are not captured. This survivor bias means the analysis conditions on successful backsliding, potentially overstating the prediction gap.

10 Conclusion

Leader traits are valuable but fundamentally limited predictors of democratic backsliding. They correctly identify high-risk leaders with high in-sample sensitivity, but they underpredict the magnitude of actual institutional erosion by a factor of $14\times$. The gap between prediction and reality provides an upper-bound estimate of the role of leader-specific factors beyond observable traits in reshaping institutions, though it also captures omitted contextual variables and model limitations.

The Venezuelan case provides a vivid illustration: once the age difference between Caldera and Chávez is accounted for, Chávez’s observable traits predict essentially no net decline in judicial independence. The entire collapse that followed the 1999 constitutional reform falls into the prediction gap, reflecting his active choices (rewriting the constitution, packing the courts, and purging judges), amplified by contextual factors such as oil wealth, political polarization, and institutional fragility.

The implication for democratic governance is sobering. We can identify who *might* backslide, but we cannot predict how far they will go. The defense against democratic erosion is therefore unlikely to succeed if it relies solely on screening leaders; it must also rely on institutional safeguards strong enough to constrain even the most determined authoritarians.

References

- Abadie, A., Diamond, A., and Hainmueller, J. (2010). Synthetic control methods for comparative case studies. *Journal of the American Statistical Association*, 105(490):493–505.
- Acemoglu, D., Naidu, S., Restrepo, P., and Robinson, J. A. (2019). Democracy does cause growth. *Journal of Political Economy*, 127(1):47–100.
- Arkhangelsky, D., Athey, S., Hirshberg, D. A., Imbens, G. W., and Wager, S. (2021). Synthetic difference-in-differences. *American Economic Review*, 111(12):4088–4118.
- Beck, N., King, G., and Zeng, L. (2000). Improving quantitative studies of international conflict: A conjecture. *American Political Science Review*, 94(1):21–35.
- Baturo, A. and Elkins, J. A. (2014). What countries select more experienced leaders? The PolEx measure of political experience. *British Journal of Political Science*, 46(2):405–425.
- Barros, R. (2002). *Constitutionalism and Dictatorship: Pinochet, the Junta, and the 1980 Constitution*. Cambridge University Press.
- Benasaglio Berlucchi, A. and Kellam, M. (2023). Who’s to blame? Populists, presidents, and dominant executives in democratic backsliding. *Democratization*, 30(5):815–835.
- Bermeo, N. (2016). On democratic backsliding. *Journal of Democracy*, 27(1):5–19.
- Besley, T. and Reynal-Querol, M. (2011). Do democracies select more educated leaders? *American Political Science Review*, 105(3):552–566.
- Coppedge, M., Gerring, J., Knutsen, C. H., Lindberg, S. I., Teorell, J., et al. (2024). V-Dem codebook v16. *Varieties of Democracy (V-Dem) Project*.
- Dreher, A., Lamla, M. J., Lein, S. M., and Somogyi, F. (2009). The impact of political leaders’ profession and education on reforms. *Journal of Comparative Economics*, 37(1):169–193.

- Funke, M., Schularick, M., and Trebesch, C. (2023). Populist leaders and the economy. *American Economic Review*, 113(12):3249–3288.
- Garoupa, N., Rosales, L., and Spruk, R. (2025). Assaults on judicial independence under the pretense of modernization: Evidence from Venezuela. Working paper, arXiv:2511.10681.
- Gandhi, J. and Przeworski, A. (2007). Authoritarian institutions and the survival of autocrats. *Comparative Political Studies*, 40(11):1279–1301.
- Geddes, B., Wright, J., and Frantz, E. (2014). Autocratic breakdown and regime transitions: A new data set. *Perspectives on Politics*, 12(2):313–331.
- Ginsburg, T. (2003). *Judicial Review in New Democracies: Constitutional Courts in Asian Cases*. Cambridge University Press.
- Ginsburg, T. and Huq, A. Z. (2018). *How to Save a Constitutional Democracy*. University of Chicago Press.
- Goemans, H. E., Gleditsch, K. S., and Chiozza, G. (2009). Introducing Archigos: A dataset of political leaders. *Journal of Peace Research*, 46(2):269–283.
- Goldstone, J. A., Bates, R. H., Epstein, D. L., Gurr, T. R., Lustik, M. B., Marshall, M. G., Ulfelder, J., and Woodward, M. (2010). A global model for forecasting political instability. *American Journal of Political Science*, 54(1):190–208.
- Graham, M. H. and Svolik, M. W. (2020). Democracy in America? Partisanship, polarization, and the robustness of support for democracy in the United States. *American Political Science Review*, 114(2):392–409.
- Greene, W. (2004). The behaviour of the maximum likelihood estimator of limited dependent variable models in the presence of fixed effects. *Econometrics Journal*, 7(1):98–119.
- Guriev, S. and Treisman, D. (2019). Informational autocrats. *Journal of Economic Perspectives*, 33(4):100–127.

- Guriev, S. and Treisman, D. (2022). *Spin Dictators: The Changing Face of Tyranny in the 21st Century*. Princeton University Press.
- Gutmann, J., Metelska-Szaniawska, K., and Voigt, S. (2024). Leader characteristics and constitutional compliance. *European Journal of Political Economy*, 82:102389.
- Haggard, S. and Tiede, L. (2024). Judicial backsliding: A guide to collapsing the separation of powers. *Democratization*, 32(2).
- Haggard, S. and Kaufman, R. R. (2021). *Backsliding: Democratic Regress in the Contemporary World*. Cambridge University Press.
- Helmke, G. (2002). The logic of strategic defection: Court–executive relations in Argentina. *American Political Science Review*, 96(2):291–303.
- Horowitz, M. C., Stam, A. C., and Ellis, C. M. (2015). *Why Leaders Fight*. Cambridge University Press.
- King, G. and Zeng, L. (2001). Logistic regression in rare events data. *Political Analysis*, 9(2):137–163.
- Jones, B. F. and Olken, B. A. (2005). Do leaders matter? National leadership and growth since World War II. *Quarterly Journal of Economics*, 120(3):835–864.
- Levitsky, S. and Ziblatt, D. (2018). *How Democracies Die*. Crown.
- Lührmann, A. and Lindberg, S. I. (2019). A third wave of autocratization is here: What is new about it? *Democratization*, 26(7):1095–1113.
- Mukand, S. W. and Rodrik, D. (2020). The political economy of liberal democracy. *Economic Journal*, 130(627):765–792.
- Mueller, H. and Rauh, C. (2022). The hard problem of prediction for conflict prevention. *Journal of the European Economic Association*, 20(6):2557–2589.

- Ríos-Figueroa, J. and Staton, J. K. (2011). An evaluation of cross-national measures of judicial independence. *Journal of Law, Economics, and Organization*, 30(1):104–137.
- Stephenson, M. C. (2003). “When the devil turns ...”: The political foundations of independent judicial review. *Journal of Legal Studies*, 32(1):59–89.
- Svolik, M. W. (2012). *The Politics of Authoritarian Rule*. Cambridge University Press.
- Svolik, M. W. (2019). Polarization versus democracy. *Journal of Democracy*, 30(3):20–32.
- Voigt, S. (2015). Determinants of *de facto* judicial independence. *International Review of Law and Economics*, 42:159–172.
- Ward, M. D., Greenhill, B. D., and Bakke, K. M. (2010). The perils of policy by p-value: Predicting civil conflicts. *Journal of Peace Research*, 47(4):363–375.

Table 1: Variable Definitions

Variable	Scale	Source	Description
<i>Panel A: Outcome Variables</i>			
Judicial constraints	0–1	V-Dem v16	Extent to which judiciary constrains executive power through review and civil liberties protection (higher = more constrained)
High court independence	Mean-zero	V-Dem v16	Expert survey: whether high courts decide independently of government wishes (higher = more independent)
Executive respects constitution	Mean-zero	V-Dem v16	Whether the executive observes constitutional provisions (higher = more respect)
<i>Panel B: Leader Traits (Archigos)</i>			
Irregular entry	0/1	Archigos 4.1	Leader entered power through coups, foreign installation, or unconstitutional means
Female leader	0/1	Archigos 4.1	Leader is female
Leader age	Years	Archigos 4.1	Current age of the leader (year – birth year)
<i>Panel C: Leader Traits (LEAD)</i>			
Military service	0/1	LEAD	Leader served in the armed forces
Combat experience	0/1	LEAD	Leader participated in armed combat
Military career	0/1	LEAD	Leader’s primary career was military
Education level	0–3	LEAD	0 = no formal education, 1 = primary/secondary, 2 = university, 3 = graduate degree
Years political experience	Years	LEAD	Years in political positions before assuming national leadership
Career politician	0/1	LEAD	Politics was the leader’s primary career
<i>Panel D: Constructed Variables</i>			
Backsliding indicator (B_{it})	0/1	Constructed	Country experienced a 5-year decline in judicial constraints exceeding 0.20

Notes: Archigos covers 188 countries, 1875–2015 (Goemans et al., 2009). LEAD covers 2,000+ leaders, 1875–2004 (Horowitz et al., 2015). V-Dem v16 covers 202 countries, 1789–2024 (Coppedge et al., 2024). Analysis sample begins in 1950.

Table 2: Summary Statistics

Variable	<i>N</i>	Mean	SD	Min	Max
<i>Panel A: Outcome Variables (V-Dem)</i>					
Judicial constraints on executive	8,975	0.522	0.313	0.003	0.991
High court independence	8,975	0.093	1.473	−3.249	3.366
Executive respects constitution	8,975	0.251	1.483	−3.236	3.352
<i>Panel B: Leader Traits (Archigos, full sample)</i>					
Irregular entry to power	8,975	0.201	0.401	0	1
Female leader	8,975	0.025	0.155	0	1
Leader age	8,962	58.3	11.6	18	104
<i>Panel C: Leader Traits (LEAD subsample, 1875–2004)</i>					
Military service	1,244	0.286	0.452	0	1
Combat experience	1,239	0.190	0.393	0	1
Military career	1,262	0.197	0.398	0	1
Education level (0–3)	1,194	2.272	0.779	0	3
Years political experience	1,201	15.47	12.49	0	72
Career politician	1,262	0.365	0.482	0	1

Notes: Panel of 174 countries. Outcome variables from V-Dem v16 (Coppedge et al., 2024), covering 1789–2024. Archigos traits from Goemans et al. (2009), covering 1875–2015. LEAD traits from Horowitz et al. (2015), covering 1875–2004. The variation in *N* across LEAD variables reflects differential coverage of specific traits. The analysis sample begins in 1950. Leader age is computed as year – birth year; in the raw data, 45 observations with implausible values (> 110 or < 0) are set to missing, of which 13 survive into the analysis sample after sample restrictions, yielding $N = 8,962$ for leader age.

Table 3: Leader Traits and Judicial Constraints on the Executive

	(1) Archigos	(2) + Military	(3) Full LEAD
Irregular entry	−0.106*** (0.021)	−0.068*** (0.012)	−0.075*** (0.013)
Female	−0.001 (0.026)	−0.054** (0.024)	−0.050** (0.024)
Leader age	−0.001* (0.001)	−0.001** (0.001)	−0.002*** (0.001)
Military service		−0.031 (0.021)	−0.021 (0.021)
Combat experience		0.013 (0.017)	0.005 (0.017)
Military career		−0.043** (0.018)	−0.038* (0.021)
Education level			0.008 (0.006)
Years experience			0.0004 (0.001)
Career politician			−0.011 (0.012)
Country & Year FE	Yes	Yes	Yes
N	8,961	1,227	1,143
Adj. R^2	0.831	0.844	0.843
Within R^2	0.066	0.111	0.122

Notes: Dependent variable: judicial constraints on the executive (0–1 scale, V-Dem). OLS with country and year fixed effects. Standard errors clustered by country in parentheses. Column (1) uses the full Archigos sample (173 clusters, 1950–2015). Columns (2)–(3) add LEAD traits, restricting the sample to countries covered by LEAD through 2004 (155 clusters). Coefficients smaller than 0.001 in absolute value are reported to four decimal places. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 4: Leader Traits and High Court Independence

	(1) Archigos	(2) + Military	(3) Full LEAD
Irregular entry	−0.402*** (0.099)	−0.286*** (0.062)	−0.354*** (0.066)
Female	0.056 (0.128)	−0.274* (0.154)	−0.249* (0.146)
Leader age	−0.006** (0.003)	−0.006** (0.003)	−0.007** (0.003)
Military service		−0.017 (0.120)	0.033 (0.122)
Combat experience		−0.041 (0.100)	−0.072 (0.098)
Military career		−0.217** (0.098)	−0.167 (0.112)
Education level			0.087** (0.034)
Years experience			−0.001 (0.002)
Career politician			0.020 (0.068)
Country & Year FE	Yes	Yes	Yes
N	8,961	1,227	1,143
Adj. R^2	0.804	0.807	0.812
Within R^2	0.042	0.070	0.092

Notes: Dependent variable: high court independence (expert survey, mean-zero, V-Dem). This variable has a different scale from the judicial constraints index in Table 3; coefficients are not directly comparable across tables. OLS with country and year fixed effects. Standard errors clustered by country in parentheses. Column (1) uses the full Archigos sample (173 clusters, 1950–2015). Columns (2)–(3) add LEAD traits, restricting the sample to countries covered by LEAD through 2004 (155 clusters). * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 5: Leader Traits and Executive Respect for the Constitution

	(1) Archigos	(2) + Military	(3) Full LEAD
Irregular entry	−0.652*** (0.116)	−0.347*** (0.088)	−0.355*** (0.095)
Female	0.070 (0.108)	−0.085 (0.137)	−0.076 (0.129)
Leader age	−0.002 (0.003)	−0.004 (0.002)	−0.007*** (0.003)
Military service		−0.304*** (0.109)	−0.291** (0.112)
Combat experience		0.139 (0.104)	0.124 (0.109)
Military career		−0.154 (0.110)	−0.110 (0.120)
Education level			0.024 (0.033)
Years experience			0.007** (0.003)
Career politician			−0.170*** (0.064)
Country & Year FE	Yes	Yes	Yes
N	8,961	1,227	1,143
Adj. R^2	0.788	0.795	0.797
Within R^2	0.082	0.091	0.101

Notes: Dependent variable: executive respect for the constitution (expert survey, mean-zero, V-Dem). This variable has a different scale from the judicial constraints index in Table 3; coefficients are not directly comparable across tables. OLS with country and year fixed effects. Standard errors clustered by country in parentheses. Column (1) uses the full Archigos sample (173 clusters, 1950–2015). Columns (2)–(3) add LEAD traits, restricting the sample to countries covered by LEAD through 2004 (155 clusters). * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 6: Leader Traits and Judicial Independence by Entry Type

	Jud. constraints		High court indep.		Exec. respects const.	
	Regular	Irregular	Regular	Irregular	Regular	Irregular
Female	−0.039 (0.025)	—	−0.207 (0.155)	—	−0.033 (0.135)	—
Leader age	−0.002** (0.001)	−0.001 (0.001)	−0.005* (0.003)	−0.005 (0.007)	−0.004 (0.003)	−0.010 (0.012)
Military service	−0.021 (0.026)	−0.049 (0.047)	0.018 (0.150)	−0.231 (0.239)	−0.233* (0.132)	−0.260 (0.347)
Combat experience	0.005 (0.026)	0.049 (0.044)	−0.058 (0.136)	0.261 (0.196)	0.039 (0.152)	0.387 (0.324)
Military career	−0.045* (0.025)	0.004 (0.049)	−0.195 (0.124)	0.244 (0.182)	−0.206 (0.156)	0.213 (0.540)
Education level	0.011* (0.007)	−0.038* (0.020)	0.099** (0.039)	−0.125 (0.112)	0.036 (0.036)	−0.211 (0.248)
Years experience	0.000 (0.001)	0.006*** (0.002)	−0.003 (0.003)	0.037*** (0.012)	0.005** (0.002)	0.042** (0.016)
Career politician	−0.009 (0.012)	−0.087* (0.045)	0.015 (0.071)	−0.351 (0.232)	−0.138** (0.066)	−0.312 (0.378)
N	920	157	920	157	920	157
Adj. R^2	0.836	0.707	0.812	0.702	0.812	0.400
Within R^2	0.018	0.088	0.067	0.138	0.072	0.103
<i>Joint F-test: all interactions = 0 (fully interacted model, $N = 1,143$)</i>						
F -statistic	0.70		1.03		1.31	
p -value	0.692		0.417		0.241	

Notes: Each column reports coefficients from $Y_{it} = \alpha_i + \gamma_t + \mathbf{X}_{it}'\beta + \varepsilon_{it}$ estimated separately on the regular-entry (elected leaders) and irregular-entry (coups/revolts) subsamples. All specifications include country and year fixed effects with standard errors clustered by country. Female is omitted in the irregular-entry subsample due to collinearity with fixed effects (near-zero variation). The three dependent variables (judicial constraints, high court independence, executive respect for the constitution) are on different scales; coefficients are not directly comparable across outcome pairs. The bottom panel reports joint F -tests from a fully interacted model on the combined LEAD sample ($N = 1,143$), testing the null that all irregular entry $\times$ trait interactions equal zero.

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 7: Predicted vs. Actual Declines During Backsliding Episodes

	Mean	SD	<i>N</i>
<i>Panel A: Baseline</i>			
Actual 5-year decline	−0.361	0.118	182
Predicted decline (from traits)	−0.026	0.075	182
Underprediction ratio		14.0×	
<i>Panel B: Sensitivity to Outliers (Trimmed Means)</i>			
Trim 5%: Actual	−0.354	0.098	164
Trim 5%: Predicted	−0.023	0.076	164
Trim 10%: Actual	−0.347	0.086	149
Trim 10%: Predicted	−0.024	0.078	149
Trim 25%: Actual	−0.347	0.056	91
Trim 25%: Predicted	−0.029	0.077	91
Trimmed ratios (5%/10%/25%)	15.5×	/ 14.5×	/ 12.0×

Notes: Panel A: 182 backsliding country-years with non-missing Archigos predictions, drawn from 51 episodes (5-year decline in judicial constraints > 0.20). Predictions from Archigos model (Column 1 of Table 3) with country and year fixed effects. The underprediction ratio is the ratio of mean actual decline to mean predicted decline. Panel B: trimmed means drop the most extreme $k\%$ of observations from both tails of the actual decline distribution and recompute the ratio. The ratio ranges from 12.0× to 15.5× across trimming levels, confirming that the baseline finding is not driven by outlier episodes.

Table 8: Binary Prediction of Backsliding: Linear Probability Model

	(1) Archigos	(2) Full LEAD
Irregular entry	0.033*** (0.012)	0.023 (0.025)
Female	−0.016* (0.009)	0.027 (0.036)
Leader age	−0.001** (0.0004)	0.0002 (0.001)
Military service		0.002 (0.017)
Combat experience		−0.049** (0.024)
Military career		0.029 (0.020)
Education level		−0.002 (0.005)
Years experience		−0.0003 (0.001)
Career politician		0.015 (0.012)
Country & Year FE	Yes	Yes
N	8,101	1,012
Adj. R^2	0.099	0.211
Within R^2	0.008	0.015
<i>In-sample classification (LEAD subsample, threshold = 0.05):</i>		
Sensitivity	100% (28/28)	
Specificity	82% (805/984)	

Notes: Dependent variable: $B_{it} = 1$ if the country experienced a 5-year decline in judicial constraints exceeding 0.20. OLS linear probability model with country and year fixed effects. Standard errors clustered by country in parentheses. The sample is restricted to country-years where the 5-year change can be computed, which requires dropping the first 5 years of observations per country; this reduces the Archigos sample from 8,961 to 8,101 and the LEAD subsample from 1,143 to 1,012. Column (1) uses the full Archigos sample; Column (2) adds LEAD traits. Classification metrics are computed in-sample from Column (2). The low threshold (0.05) is chosen to maximize sensitivity in an early warning application. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 9: Rare Events Logit: Predicting Backsliding Episodes

	Archigos		Full LEAD	
	(1) Logit	(2) Relogit	(3) Logit	(4) Relogit
Irregular entry	1.316*** (0.152)	1.314*** (0.152)	0.739 (0.473)	0.735 (0.467)
Female	-1.094 (1.012)	-0.605 (1.011)	0.443 (1.063)	0.840 (1.050)
Leader age	-0.029*** (0.007)	-0.029*** (0.007)	-0.013 (0.024)	-0.011 (0.023)
Combat experience			-0.689 (0.524)	-0.642 (0.518)
Military career			0.794 (0.510)	0.785 (0.504)
Education level			-0.012 (0.225)	-0.035 (0.222)
Years experience			-0.044* (0.025)	-0.041* (0.025)
Career politician			1.553*** (0.489)	1.491*** (0.483)
Decade FE	Yes	Yes	Yes	Yes
N	8,080	8,080	1,031	1,031
<i>Classification (threshold = 0.05):</i>				
Sensitivity	40.4%	40.4%	42.9%	46.4%
Specificity	90.1%	89.9%	87.0%	81.8%

Notes: Dependent variable: $B_{it} = 1$ if the country experienced a 5-year decline in judicial constraints exceeding 0.20. Columns (1) and (3): standard logistic regression. Columns (2) and (4): rare events logit following King and Zeng (2001). All models include decade fixed effects (2010s as reference category) instead of country and year fixed effects, because the rare events logit estimator cannot accommodate high-dimensional fixed effects. Columns (1)–(2) use the full Archigos sample with only Archigos-coded traits (irregular entry, female, leader age); Columns (3)–(4) add LEAD traits on the smaller LEAD subsample. Military service is omitted from Columns (3)–(4) due to near-perfect collinearity with military career in the logit specification. Standard errors (robust) in parentheses. Classification metrics computed at a 0.05 probability threshold. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 10: Venezuela: Synthetic Difference-in-Differences Estimates

Specification		ATT	SE	<i>t</i> -stat
<i>Panel A: Main Estimates (1985–2015)</i>				
(1)	Difference-in-differences	−0.650***	0.179	−3.62
(2)	Synthetic control	−0.599***	0.142	−4.22
(3)	Synthetic DiD (preferred)	−0.633***	0.215	−2.94
<i>Panel B: Leave-One-Out Robustness</i>				
(4)	Excluding Nicaragua & Ecuador	−0.659***	0.255	−2.58
(5)	Excluding Colombia	−0.634***	0.176	−3.60
(6)	Excluding Chile	−0.626***	0.161	−3.89
(7)	Excluding Brazil	−0.639***	0.166	−3.85
<i>Panel C: Placebo Test</i>				
(8)	In-time placebo (1992)	−0.011	0.158	−0.07
<i>Panel D: Extended Pre-Treatment Window</i>				
(9)	SDID, 1960–2015	−0.698***	0.198	−3.53
Donor pool: 17 Latin American countries. Main period: 1985–2015.				

Notes: Dependent variable: judicial constraints on the executive (0–1 scale, V-Dem). Treatment: Venezuela, 1999 Bolivarian Constitution. Panels A–C use the 1985–2015 window; Panel D extends the pre-treatment period to 1960 as a robustness check. Panel A compares three estimation methods. Panel B drops individual donors from the SDID specification. Panel C assigns a false treatment date of 1992 using only pre-1999 data. Panel D re-estimates SDID with the full 1960–2015 window. Standard errors from placebo inference (50 replications), which assigns treatment to each donor in turn. The small donor pool (17 countries) means placebo p-values are approximate. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 11: Leader Trait Changes: Caldera to Chávez

Trait	Caldera	Chávez	Δ	$\hat{\beta} \times \Delta$
Irregular entry	0	0	0	0.000
Female	0	0	0	0.000
Leader age	78	45	-33	+0.066
Military service	0	1	+1	-0.021
Combat experience	0	1	+1	+0.005
Military career	0	1	+1	-0.038
Education level	3	2	-1	-0.008
Years experience	30	5	-25	-0.010
Career politician	1	0	-1	+0.011
Trait-predicted decline				+0.005
Total SDID estimate				-0.633
Prediction gap				-0.638

Notes: Coefficients ($\hat{\beta}$) from Column (3) of Table 3. The column $\hat{\beta} \times \Delta$ shows the predicted change in judicial independence attributable to each trait change from Caldera to Chávez. Caldera: civilian career politician, law degree, 30+ years of experience, age 78 at transition. Chávez: military officer, 1992 coup participant, political newcomer, age 45. Chávez's entry is coded as "regular" in Archigos because he won the 1998 election, so irregular entry contributes zero. The large age difference (-33 years) produces a positive predicted effect (+0.066) that offsets the negative military and experience contributions, yielding a net trait-predicted change near zero. Individual contributions may not sum exactly to the total due to rounding of displayed coefficients; the total is computed from full-precision estimates. This is a descriptive accounting exercise using correlational panel coefficients; the prediction gap captures omitted variables and model misspecification in addition to leader-specific factors.

Table 12: Robustness of the Prediction Gap to Alternative Backsliding Definitions

Definition	Episodes	Country-yrs	Avg. actual	Avg. predicted	Ratio
CI non-overlap (H&T 2024)	55	197	−0.339	−0.020	16.7×
Threshold > 0.15	82	290	−0.292	−0.021	14.2×
Threshold > 0.20 (baseline)	51	184	−0.361	−0.026	14.0×
Threshold > 0.25	39	141	−0.402	−0.019	20.8×

Notes: Each row defines backsliding differently and reports the mean actual 5-year decline in judicial constraints on the executive, the mean predicted decline (from Archigos traits with country and year fixed effects), and their ratio. The CI non-overlap method follows Haggard and Tiede (2024): backsliding occurs when the upper bound of the current 95% credible interval for judicial independence falls below the lower bound in any of the preceding five years. Threshold definitions use fixed cutoffs applied to the 5-year change. “Episodes” counts distinct onset events (a new episode begins when a country enters backsliding after at least one non-backsliding year). The prediction gap is robust across all definitions, ranging from 14.0× to 20.8×. No hypothesis tests are reported; all ratios are descriptive measures.

Table 13: Prediction Gap Across Model Specifications

Model	Avg. predicted	Avg. actual	Ratio
OLS with country/year FE	−0.027	−0.361	13.6×
Quantile reg. ($\tau = 0.10$)	−0.010	−0.361	35.7×
Quantile reg. ($\tau = 0.25$)	−0.017	−0.361	21.5×
Quantile reg. ($\tau = 0.50$)	−0.017	−0.361	20.9×
Random forest (LOCO-CV)	−0.088	−0.361	4.1×
Gradient boosting (LOCO-CV)	−0.078	−0.361	4.6×

Notes: Each row estimates the relationship between leader traits (irregular entry, female, leader age) and judicial constraints on the executive using a different model specification, then computes the ratio of mean actual to mean predicted 5-year declines during backsliding episodes (5-year drop > 0.20). OLS and quantile regressions include country and year fixed effects estimated on the full Archigos sample. Random forest and gradient boosting use leave-one-country-out cross-validation (LOCO-CV) with year dummies but no country dummies, providing honest out-of-sample predictions. The prediction gap persists across all specifications, confirming that the underprediction is not an artifact of linear functional form. No hypothesis tests are reported; all ratios are descriptive measures.




Figure 1: Predicted vs. Actual 5-Year Changes in Judicial Independence

Notes: Each point is a country-year observation from the Archigos sample ($N = 8,961$, 174 countries, 1955–2015). The horizontal axis shows the predicted 5-year change in judicial constraints on the executive (from fitted values of Column 1 of Table 3 with country and year fixed effects). The vertical axis shows the actual 5-year change. Gray circles: non-backsliding observations. Red circles: backsliding episodes (5-year decline > 0.20). Dashed line: perfect prediction (45-degree line). Prediction intervals not shown; the scatter reveals the distribution of prediction errors. Backsliding episodes cluster far below the 45-degree line, indicating dramatic underprediction.




Figure 2: Actual vs. Predicted Backsliding Magnitude: Top 10 Episodes

Notes: The ten most severe backsliding episodes by magnitude of 5-year decline in judicial constraints on the executive ($N = 10$ episodes out of 51 total). Episodes shown: Belarus 1996 (Lukashenko), Philippines 1972 (Marcos), Republic of the Congo 1968 (Debat), Chile 1974 (Pinochet), Zambia 1972 (Kaunda), Cuba 1959 (Castro), Venezuela 2000 (Chávez), Comoros 1975, Peru 1993 (Fujimori), Turkey 2016 (Erdogan). Red bars: actual 5-year decline. Blue bars: predicted decline from Archigos leader traits (Column 1 of Table 3 with country and year fixed effects). Each bar represents a single episode at its onset year. In all ten cases, the actual decline substantially exceeds the prediction.




Figure 3: Venezuela vs. Synthetic Control: Judicial Independence

Notes: Synthetic difference-in-differences (Arkhangelsky et al., 2021). Outcome variable: judicial constraints on the executive (0–1 scale, V-Dem). Period: 1985–2015. Donor pool: 17 Latin American countries (Argentina, Bolivia, Brazil, Chile, Colombia, Costa Rica, Dominican Republic, Ecuador, El Salvador, Guatemala, Honduras, Mexico, Nicaragua, Panama, Paraguay, Peru, Uruguay). Solid line: Venezuela. Dashed line: synthetic Venezuela (SDID-weighted combination of donors). Vertical dashed line at 1999 marks the Bolivarian Constitution. In the pre-treatment period (1985–1998), Venezuela fluctuates between 0.65 and 0.75 while the synthetic control remains between 0.53 and 0.57; the level difference is absorbed by unit fixed effects in the SDID estimator. Post-1999, Venezuela’s judicial independence collapses while the synthetic control remains stable. See Table 10 for standard errors from placebo inference.




Figure 4: Robustness: Predicted vs. Actual Changes Using CI Non-Overlap Definition

Notes: Replication of Figure 1 using the Haggard and Tiede (2024) CI non-overlap definition of backsliding. Red circles indicate country-years where the upper bound of the current 95% credible interval for judicial independence falls below the lower bound in any of the preceding five years. Gray circles: all other observations. Dashed line: perfect prediction. Axes as in Figure 1. Archigos sample ($N = 8,961$, 174 countries, 1955–2015). The pattern of dramatic underprediction for backsliding episodes is robust to this alternative, statistically grounded definition.